
Nonstationarity Manufactures Heavy Tails: Scale-Normalized Sequential Learning

Anonymous Authors¹

Abstract

We study online prediction on streams whose loss-gradient *scale* is strongly nonstationary, using a large financial dataset as a testbed. The gradient-norm scale drifts several-fold across the record with overnight jumps, and this drift alone makes the *pooled* gradient look heavy-tailed (tail index $\hat{\alpha} \approx 2.4$); the *causally scale-normalized* gradient—the quantity a learner should actually act on—is markedly lighter ($\hat{\alpha} \approx 3.7$), so the difficulty is dominated by nonstationarity rather than by an intrinsically heavy tail. In this regime plain online gradient descent (OGD) has no usable learning rate—small rates underfit while any rate large enough to learn diverges when the scale spikes—and clipping to a *scale-dependent* threshold is fragile because the clipped step still inherits the drift. We propose *Scale-Normalized OMD* (SN-OMD), which divides the gradient by a tracked robust scale and caps it at a *scale-free* constant M ; the per-step move is then bounded by M rather than by a regime-dependent quantity, so the iterates provably remain bounded for any learning rate or realization (a stability, not convergence, guarantee). We prove a high-probability dynamic-regret bound of order $T^{1/p}$ under a finite, possibly time-varying p -th moment ($p \in (1, 2]$), and give a budget calculation—corroborated by a direct measurement—indicating that the interaction of tail drift with a moving comparator pushes *dynamic* regret well above $\sqrt{T}$. Empirically, on a continuous multi-regime stream SN-OMD stays bounded across the entire tested learning-rate range and is the most accurate of the methods compared, whereas OGD and scale-dependent clipping both diverge above $\text{lr} = 10^{-2}$.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

1. Introduction

Financial markets are a canonical source of heavy-tailed, distribution-free data: returns and microstructure signals exhibit power-law tails, occasional data errors act as adversarial contamination, and assumptions such as sub-Gaussianity are untenable. Sequential prediction on such streams therefore calls for *distribution-free* guarantees that hold under weak moment assumptions only. A rich line of work in robust statistics shows that the empirical mean is a poor estimator under heavy tails, whereas estimators such as Catoni’s M-estimator, median-of-means, and the trimmed mean achieve sub-Gaussian deviations assuming only a finite variance (Catoni, 2012; Devroye et al., 2016; Lugosi & Mendelson, 2019); in optimization, gradient clipping is the corresponding algorithmic primitive for heavy-tailed gradient noise (Zhang et al., 2020; Gorbunov et al., 2020; Nguyen et al., 2023).

This paper concerns a facet these lines mostly abstract away: on real streams the gradient *scale* is strongly nonstationary, and this—more than an intrinsically heavy tail—is what breaks existing methods. Measuring the gradient of an online linear predictor on the Jane Street market dataset, we find (i) the gradient *scale* drifts roughly $6\times$ across two months with overnight jumps; (ii) this drift alone inflates the *pooled* gradient-norm tail index to $\hat{\alpha} \approx 2.4$ —a scale mixture of lighter-tailed increments looks heavy-tailed—whereas the *causally scale-normalized* gradient $\|g_t\|/s_t$, which is what a learner actually acts on, is much lighter, $\hat{\alpha} \approx 3.7$; and (iii) the residual tail is nonetheless heavier than either its feature or residual factor alone, an *interaction* effect from their co-movement. The operative difficulty is thus nonstationarity of scale, with only a mild residual heavy tail. This regime—a drifting scale against a moving comparator—is exactly where existing distribution-free online learning results, which are either in expectation, on bounded domains, or against a static comparator (Zhang et al., 2026; Zhang & Cutkosky, 2022), do not reach.

The consequences are sharp. Plain OGD has *no* usable learning rate on such a stream: a small rate underfits, while any rate large enough to learn diverges the first time the gradient scale spikes in a turbulent regime. The textbook fix—clip the gradient at a threshold τ_t tracking the local scale—does

not help under a drifting scale, because the clipped step is $\propto \min(\|g_t\|, \tau_t)$ and so still scales *with* the regime: a turbulent window produces large steps and diverges. We show this failure both empirically and as a structural property of any scale-dependent threshold.

Our approach. The fix is to decouple the step from the tail scale. *Scale-Normalized OMD* (SN-OMD) divides the gradient by a predictable robust scale s_t and caps the result at a *scale-free* constant M ,

$$\hat{g}_t = \text{clip}(g_t/s_t, M), \quad w_{t+1} = \Pi_{\mathcal{W}}\left(w_t - \frac{\eta}{\sqrt{t}} \hat{g}_t\right).$$

The single factor $1/s_t$ replaces clipping’s *scale-dependent* step bound $M s_t$ with a *scale-free* bound M . Because $\|\hat{g}_t\| \leq M$ deterministically, the iterates remain bounded for *any* learning rate or realization—an unconditional *stability* guarantee (not, on its own, convergence) that needs no moment assumption.

Contributions.

- **Problem and a budget heuristic (Section 2).** We formalize distribution-free online learning under a *drifting* p -th moment, for which the per-regime achievable rate is $T^{1/p}$, and give a heuristic budget calculation—corroborated by a direct measurement—suggesting that the interaction of tail drift with a moving comparator pushes *dynamic* regret well above $\sqrt{T}$ while static regret stays near $\sqrt{T}$. This is a motivating observation, not a matching lower bound: a proper path-length-constrained Ω is left open.
- **Algorithm and theory (Section 3).** We propose SN-OMD and prove (Thm. 3.1) unconditional boundedness of the iterates and (Thm. 3.3) a high-probability dynamic-regret bound $\tilde{O}((\sup_t \sigma_t)(D + \sqrt{DP_T}) T^{1/p})$, separating the tail (rate exponent) from the scale drift (a constant).
- **Experiments (Section 4).** On Jane Street, SN-OMD stays bounded across the whole tested learning-rate range and attains the highest weighted R^2 , whereas OGD and scale-dependent clipping both diverge above $\text{lr} = 10^{-2}$. A cap-sweep shows the finite cap is a sweet spot between normalized-GD and uncapped scale-adaptive OGD, and we empirically (and, under a mild drift model, provably) discharge the scale-tracking assumption; a tested implementation is released.

2. Problem Setup

Protocol. Learning proceeds in rounds $t = 1, \dots, T$. At round t the learner plays $w_t \in \mathcal{W} \subseteq \mathbb{R}^d$ (a convex set of di-

ameter D ; the unbounded case is discussed in Section 3), observes a convex loss ℓ_t (for us the sample-weighted squared error $\ell_t(w) = \omega_t(\langle w, x_t \rangle - y_t)^2$), and receives a stochastic subgradient g_t with $\mathbb{E}[g_t \mid \mathcal{F}_{t-1}] = \bar{g}_t \in \partial \ell_t(w_t)$, where $(\mathcal{F}_t)$ is the natural filtration and $\varepsilon_t = g_t - \bar{g}_t$ is the noise.

Drifting heavy-tailed noise. We assume only a *finite, time-varying* low-order moment:

Assumption 2.1 (Drifting p -th moment). For each t , $\mathbb{E}[\|\varepsilon_t\|^{p_t} \mid \mathcal{F}_{t-1}] \leq \sigma_t^{p_t}$ for some $p_t \in (1, 2]$, where both the scale σ_t and the tail index p_t may drift with t , and $p_t < 2$ (infinite variance) is permitted. Write $p := \inf_t p_t$.

No sub-Gaussianity, boundedness, or stationarity is assumed. *Distribution-free* means the learner knows none of σ_t, p_t, D , or the comparator norm. The regret analysis does additionally require a *predictable* scale tracker (Assumption 3.2); we do not treat this as free, and discharge it under a mild drift model (Proposition A.5).

Dynamic regret. Performance is measured against a moving comparator $u_{1:T}$ with path length $P_T = \sum_t \|u_t - u_{t-1}\|$,

$$R_T(u_{1:T}) = \sum_{t=1}^T [\ell_t(w_t) - \ell_t(u_t)] \leq \sum_{t=1}^T \langle \bar{g}_t, w_t - u_t \rangle, \quad (1)$$

the notion appropriate for regime-switching streams; static regret is the special case $u_t \equiv u$. We report the sample-weighted zero-mean $R^2 = 1 - \sum_i \omega_i (y_i - \hat{y}_i)^2 / \sum_i \omega_i y_i^2$ on data.

2.1. A budget heuristic for the dynamic wall

The achievable rate degrades smoothly with the tail. The core primitive is robust mean estimation: under a finite p -th moment the best estimator of a mean from n samples has error $\Theta(\sigma n^{-(1-1/p)})$ (Devroye et al., 2016; Lugosi & Mendelson, 2019), which interpolates the sub-Gaussian rate ($p = 2$) and no concentration ($p \rightarrow 1$). Aggregated online this yields, and is matched by, a regret of order $T^{1/p}$ (Zhang et al., 2026):

$$R_T \asymp \sigma \cdot T^{1/p}, \quad T^{1/2} (p=2) \longrightarrow T (p \rightarrow 1). \quad (2)$$

Two refinements matter under nonstationarity. Partitioning $[T]$ into regimes r of length L_r , index p_r , and scale σ_r , and applying the bound within each, *static* regret (one comparator) is governed by the worst regime, while *dynamic* regret (a comparator per regime) has the per-regime costs *add*:

$$R_T^{\text{stat}} \gtrsim \max_r \sigma_r L_r^{1/p_r}, \quad R_T^{\text{dyn}} \gtrsim \sum_r \sigma_r L_r^{1/p_r}. \quad (3)$$

Because $L^{1/p}$ is convex, heavy-tail cost cannot be averaged across regimes: a few short heavy regimes ($p_r < 2$)

Algorithm 1 Scale-Normalized OMD (SN-OMD)

Input: cap M , step η , predictable scale tracker $\mathcal{S}$; $w_1 = 0$.

for $t = 1$ **to** T **do**

play w_t ; observe ℓ_t and stochastic subgradient g_t

$s_t \leftarrow \mathcal{S}.\text{scale}$ (predictable, uses $g_{t'}$ for $t' < t$)

$\hat{g}_t \leftarrow \text{clip}(g_t/s_t, M)$

$w_{t+1} \leftarrow \Pi_{\mathcal{W}}(w_t - \frac{\eta}{\sqrt{t}} \hat{g}_t)$

$\mathcal{S}.\text{update}(\|g_t\|)$

end for

dominate the sum. Two caveats keep (3) honest. First, it is a per-regime *budget*, not a proved regret lower bound at a fixed path length: it charges each regime its worst-case cost against a *fresh* comparator, so it lower-bounds dynamic regret only against an *unconstrained* comparator sequence—which is trivially $\Omega(T)$ in any case. The interesting object is regret at a given path length P_T , which (3) does not constrain. Second, it separates drift from the static case only qualitatively: prior heavy-tailed results fix a *single* p , a stationary scale, and a *static* comparator, and recover the $T^{1/p}$ rate (2) (Zhang et al., 2026; Zhang & Cutkosky, 2022; Bubeck et al., 2013), whereas here the per-regime costs *add* across a drifting exponent, and the static budget (a max, not a sum) does not see this. Empirically the dynamic budget does grow super- $\sqrt{T}$: measured on Jane Street (Section 4), its ratio to $\sqrt{T}$ *increases* with horizon (from ≈ 1.3 to ≈ 5.9 ; fitted exponent ≈ 0.9), while the static budget stays $\approx 0.16\sqrt{T}$. We therefore read (3) as a *motivating heuristic*—the combination of tail drift and a moving comparator makes near- $\sqrt{T}$ *dynamic* regret implausible on this data—rather than as a matching lower bound; a proper P_T -constrained Ω is left open. What SN-OMD does deliver is to attain the per-regime worst-tail rate $T^{1/p}$ rather than diverge.

3. Scale-Normalized Online Mirror Descent

Algorithm. SN-OMD maintains a *predictable* robust estimate $s_t > 0$ of the gradient-norm scale (an $\mathcal{F}_{t-1}$ -measurable function of past gradient norms) and takes a scale-invariant, constant-capped step (Algorithm 1). Writing $v_t = g_t/s_t$ and $\text{clip}(v, M) = v \cdot \min\{1, M/\|v\|\}$, the update is $w_{t+1} = \Pi_{\mathcal{W}}(w_t - \frac{\eta}{\sqrt{t}} \text{clip}(v_t, M))$. The scale s_t can be any predictable robust tracker; we use a winsorized exponential moving average of $\|g_t\|$ by default, and analyze a two-timescale “peak-hold” variant below.

Why normalize, not clip. Clipping to a scale-dependent threshold $\tau_t = c s_t$ gives a step $\propto \min\{\|g_t\|, c s_t\}$, which scales *with* the drifting s_t : a turbulent regime yields large steps and divergence. Dividing by s_t first gives a step $\propto$

$\min\{\|g_t\|/s_t, M\}$, bounded by the *scale-free* constant M . The two differ by the factor $1/s_t$, and this single factor is what removes the scale-drift transmission.

Theorem 3.1 (Stability, unconditional). *For any inputs, $\|\hat{g}_t\| \leq M$ deterministically. Hence on a bounded domain $\|w_t\| \leq D$ for all t , and on $\mathbb{R}^d$ (step $\eta/\sqrt{t}$) $\|w_t\| \leq \|w_1\| + 2M\eta\sqrt{t}$. Hence the iterates remain uniformly bounded (on $\mathcal{W}$) or grow at most as $O(\sqrt{t})$ (on $\mathbb{R}^d$) for any learning rate, scale drift, or realization—a stability guarantee, not by itself convergence, that uses no moment assumption.*

This is the crucial contrast with scale-dependent clipping, whose per-step move $\propto s_t$ is bounded only by the drifting scale.

Regret. The scale tracker must track the local scale. We assume:

Assumption 3.2 (Scale tracking). The predictable scale brackets the local noise scale: $c_1\sigma_t \leq s_t \leq c_2\sigma_t$ for constants $0 < c_1 \leq c_2$.

Only the lower bracket is essential—it bounds $\|v_t\|$ and hence the rate, and under-estimating the scale is the sole divergence risk—while over-estimation merely shrinks the step η/s_t into mild underfitting, never divergence. Assumption 3.2 is thus *one-sided* in the direction that matters, which is what makes it dischargeable rather than an oracle: under a mild piecewise-slowly-varying drift model, a two-timescale tracker *provably* meets the lower bracket off an $O(NW)$ -round set (Proposition A.5), and we verify in Section 4.3 that it does so at *every* round on the real gradient-scale process while keeping its upward variation to a small constant multiple of the true scale drift. Define that upward variation $W_s = s_1 + \sum_t (s_t - s_{t-1})_+$ and the scale-weighted path $P_T^s = \sum_t s_t \|u_t - u_{t-1}\|$.

Theorem 3.3 (Dynamic regret). *Under Assumptions 2.1 and 3.2, SN-OMD with $M = \Theta(T^{1/p})$ (or a doubling schedule, requiring no knowledge of p) satisfies, with probability $\geq 1 - \delta$, simultaneously for all comparators $u_{1:T}$,*

$$\begin{aligned} R_T(u_{1:T}) &= O\left((D\sqrt{W_s} + \sqrt{DP_T^s})\sqrt{S_1} T^{(2-p)/(2p)}\right) \\ &= \tilde{O}\left((\sup_t \sigma_t) (D + \sqrt{DP_T}) T^{1/p} \log(1/\delta)\right), \end{aligned} \quad (4)$$

where $S_1 = \sum_t \sigma_t$; the second line uses $S_1 \leq T \sup_t \sigma_t$ and $W_s, P_T^s \lesssim \sup_t \sigma_t (1, P_T)$ (so it is tight up to the scale normalization $\sup_t \sigma_t \gtrsim 1$). Together with Theorem 3.1, the iterates are bounded for every learning rate and realization.

Proof sketch. Analyze SN-OMD in its *adaptive-step* form: it is OGD with step $\eta_t = \eta/s_t$ on the clipped raw gradient. The one-step mirror-descent inequality, summed with a summation-by-parts over the varying step, telescopes to

$\sum_t s_t (\|w_t - u_t\|^2 - \|w_{t+1} - u_t\|^2) \leq D^2 W_s + 2D P_T^s$ —the only place the two nonstationarities enter. The curvature term is controlled by the clipped second moment $\|\text{clip}(v, M)\|^2 \leq M^{2-p} \|v\|^p$; the clipping bias by $\|v - \text{clip}(v, M)\| \leq \|v\|^p / M^{p-1}$; and the martingale part by Freedman’s inequality (high probability). Optimizing η and $M = \Theta(T^{1/p})$ balances curvature against bias and gives (4). Full proofs, including per-round p_t and the unknown- p /unknown- $\|u\|$ reductions, are in the appendix. $\square$

Remark 3.4 (Tail vs. scale drift). The bound *separates* the two effects: the tail index p sets the *rate* exponent (through $M^{2-p} = T^{(2-p)/p}$, giving the overall $T^{1/p}$), while the scale drift enters only through the factor $\sqrt{W_s}$ —a multiplicative *constant* on that rate, not a change to the exponent. Crucially, dividing by s_t does not smuggle a $\sup_t \sigma_t$ *rate* penalty; the drift is a bounded constant whenever W_s is (Remark 3.5 and Proposition A.5).

Remark 3.5 (Tracker design). The bound degrades with the tracker’s upward variation W_s . A peak-hold envelope $s_t = \max\{c \cdot \text{median}(\text{recent } \|g_t\|), (1-\rho)s_{t-1}\}$ has $W_s \approx \sup_t s_t + \rho \sum_t s_t$; choosing $\rho \asymp 1/L$ (the regime length) gives $W_s = O(V_\sigma^+)$ and the dynamic-optimal rate, whereas too large a ρ churns and makes $W_s = \Theta(T)$. An unknown regime length is handled by a strongly-adaptive wrapper (Daniely et al., 2015; Cutkosky, 2020) using SN-OMD’s interval-regret property (Theorem 3.1 gives the required bounded feedback).

4. Experiments

Setup. We use the Jane Street Real-Time Market Data Forecasting dataset (Jane Street, 2024): 47,127,338 rows over 1,699 trading days, 79 features, 39 symbols; target `responder_6`; metric the sample-weighted zero-mean R^2 . We stream chronologically and compare plain OGD, quantile-based AdaptiveClip, RobustOMD (a scale-dependent robust-mean clip), and SN-OMD. **Protocol caveat.** Features here are standardized using full-sample per-feature statistics over each evaluated slice; because this peeks at the whole slice it is a look-ahead confound that *inflates* absolute R^2 . The values below are therefore *not* comparable to the competition’s (much smaller) held-out leaderboard score, and are reported only for *cross-method* comparison—every method shares the identical preprocessing—and for the stability contrast, which is unaffected. We additionally report a *causal* (running-statistics) re-run that removes this look-ahead. All tail measurements are learner-independent (gradients at the fixed best linear predictor w^*).

The scale is nonstationary; the pooled tail is mostly a scale mixture. At the fixed optimum w^* , the *pooled* gradient-norm tail index is $\hat{\alpha} \approx 2.43$, and it is an *inter-*

action effect: the residual and feature norms each have $\hat{\alpha} \approx 4$, but their product—the gradient—has a heavier tail because extreme features and residuals co-occur (Figure 1a). Most of that pooled heaviness, however, is manufactured by nonstationarity. The daily gradient scale drifts $\approx 6\times$ with overnight jumps (Figure 1b), and once we divide by a *causal* robust scale s_t —the quantity SN-OMD actually acts on—the tail lightens markedly, to $\hat{\alpha} \approx 3.7$ (winsorized-EMA s_t) or ≈ 2.9 (trailing median), i.e. toward finite higher moments. So the operative difficulty is the scale drift, leaving only a mild residual heavy tail. (Per-day $\hat{\alpha}$ dips below 2 on a minority of days, but single-day Hill estimates are noisy and we do not rely on them.) As the heuristic of Section 2.1, the *dynamic* budget $\sum_r \sigma_r L_r^{1/p_r}$ grows super- $\sqrt{T}$ —its ratio to $\sqrt{T}$ rises from ≈ 1.3 to ≈ 5.9 across horizons (fitted exponent ≈ 0.9)—while the *static* budget stays $\approx 0.16\sqrt{T}$; we read this as motivating, not as a proved bound.

OGD and scale-dependent clipping both diverge; only normalization is bounded. We stream one learner through the multi-regime window and sweep its learning rate under the causal protocol (Figure 3, Table 1). Plain OGD has no usable rate: below $\text{lr} \approx 10^{-2}$ it is stable but underfits, and at any larger rate it diverges the first time the gradient scale spikes. Crucially, *scale-dependent* robust clipping does no better—AdaptiveClip and RobustOMD (MoM) diverge at the *same* threshold ($\text{lr} \geq 2 \times 10^{-2}$) and reach essentially the same best R^2 as OGD—because a threshold $\tau_t = c s_t$ rescales *with* the regime, transmitting the $6\times$ scale swing into the step (Figure 2). Truncating relative to the current scale guards against spikes but not against a regime-wide scale increase. (An earlier version of this comparison showed a large clipping advantage; that was an artifact of full-sample feature standardization and vanishes under the causal protocol of the Setup—a cautionary point in its own right.)

Scale normalization is bounded and most accurate. SN-OMD divides by a *predictable* tracked scale s_t and caps the result at a scale-free M , so its step is bounded by M regardless of the regime. It stays bounded across the *entire* learning-rate range (peak rolling loss ≈ 5 throughout) while OGD and both clippers diverge, and it attains the highest weighted R^2 (Table 1). The mechanism is isolated in Section 4.1: the gain over the divergent baselines comes from *scale normalization* (dividing by s_t , which preserves gradient magnitude within a regime), and the cap is what keeps it safe when the tail is genuinely heavy.

4.1. The cap M interpolates normalized-GD and scale-adaptive OGD

SN-OMD is a one-parameter family in the cap M . Whenever $\|g_t\|/s_t > M$, $\text{clip}(g_t/s_t, M) = M g_t/\|g_t\|$ and s_t




Figure 1. The scale drifts; the pooled tail is largely a scale mixture. (a) The *pooled* gradient tail (Hill index $\hat{\alpha} \approx 2.43$) is heavier than either its residual or feature factor ($\hat{\alpha} \approx 4$)—an interaction effect from co-movement. (b) The daily gradient tail index across the full 1,699-day record drifts; per-day estimates occasionally fall below 2 but single-day Hill estimates are noisy. Causally normalizing by the scale lightens the tail to $\hat{\alpha} \approx 3.7$, so the pooled heaviness is largely nonstationarity (Section 4).

Table 1. Best weighted R^2 (each method at its own tuned rate) and stable learning-rate range on the continuous Jane Street stream, *causal* standardization. Scale-dependent methods (top) all diverge above $\text{lr} = 10^{-2}$ at essentially the same R^2 ; scale-invariant methods (bottom) stay bounded to the largest rate tested, and among them magnitude-preserving caps beat the magnitude-blind normalized-GD endpoint. SN-OMD’s cap ($M=5$) is best. *Absolute* R^2 is not comparable to the competition’s held-out score; the cross-method comparison is fair (shared preprocessing).

Method	Best weighted R^2	Stable lr
<i>Scale-dependent (diverge above 10^{-2}):</i>		
OGD	0.018	$\leq 10^{-2}$
AdaptiveClip	0.013	$\leq 10^{-2}$
RobustOMD (MoM)	0.014	$\leq 10^{-2}$
<i>Scale-invariant (bounded):</i>		
Normalized-GD ($M \rightarrow 0$)	0.087	$\leq 5 \times 10^{-1}$
Scale-adaptive OGD ($M \rightarrow \infty$)	0.162	$\leq 5 \times 10^{-1}$
SN-OMD ($M=5$, ours)	0.169	$\leq 5 \times 10^{-1}$

cancels, so $M \rightarrow 0$ recovers *normalized-GD* $w_{t+1} = w_t - (\eta M / \sqrt{t}) g_t / \|g_t\|$, while $M \rightarrow \infty$ gives *scale-adaptive OGD* $w_{t+1} = w_t - (\eta / \sqrt{t}) g_t / s_t$ with no truncation. Both endpoints satisfy Theorem 3.1 (normalized-GD trivially, as $\|g_t\| / \|g_t\| = 1$), so non-divergence is *not* what singles out SN-OMD. The distinction is that a finite M *preserves gradient magnitude within the bulk* of the distribution—the step shrinks as the learner nears a good predictor—and truncates only the tail, whereas normalized-GD discards magnitude on every round. The tracked scale s_t is what makes “the bulk” regime-independent: the membership test $\|g_t\| / s_t \leq M$ is scale-free, which a fixed threshold cannot be under $6 \times$ drift.

Sweeping M (each cap at its own tuned step) confirms an interior optimum. On a synthetic heavy-tailed stream ($p = 1.5$, static reachable comparator; Figure 4) the best cap sits strictly between the endpoints—about $240 \times$ better than the uncapped end (heavy-tail instability) and $10 \times$ better than the normalized-GD end (magnitude-blind oscillation). On the Jane stream, where the *normalized* tail is only mildly heavy ($\hat{\alpha} \approx 3.7$), the picture is consistent but flatter: normalized-GD ($M \rightarrow 0$) is clearly worst while a broad plateau of finite M matches the uncapped end (Table 1). So on real data the win is magnitude preservation, and the cap is a low-cost safeguard that becomes essential only when the tail is genuinely heavy—exactly the regime the synthetic isolates.

4.2. Synthetic streams with known (σ_t, p_t)

To test the theory’s mechanism with a known ground-truth tail, we generate online regression streams with controlled (σ_t, p_t) (Student- t noise of index p ; Figure 5). As p decreases from 2 to 1.3, plain OGD’s fitted cumulative-regret exponent rises $0.37 \rightarrow 0.70 \rightarrow 1.55$ —super-linear (divergent) once the tail is heavy—while SN-OMD stays sublinear. This reproduces the real-data divergence with a *known* tail index, showing the tail is the mechanism, not a learning-rate artifact. SN-OMD’s measured exponents sit *below* $1/p$: the stochastic stationary instance is easier than the worst case, so this is consistent with, but does not tighten, Theorem 3.3 (a tight verification would need the adversarial lower-bound instance, which we do not have). Against a drifting comparator, SN-OMD’s dynamic regret grows with




Figure 2. **Why scale-dependent clipping fails.** (a) The gradient scale drifts $\sim 6\times$ across days. (b) A fixed-window adaptive clip threshold cannot win: a short window tracks the scale (so the step inherits the swing), a long window is smooth but lags regime onsets.

the path length at a fitted exponent ≈ 0.43 , consistent with the $\sqrt{DP_T}$ term.

4.3. Discharging the scale-tracking assumption

The regret bound (Theorem 3.3) rests on Assumption 3.2, that a *predictable* scale s_t brackets the local noise scale. We close this loop empirically on the real gradient-scale process (357,000 gradient norms at w^* , ground-truth scale a centered rolling median, all trackers predictable), where the true upward variation is V_σ^+ . The relevant requirement is one-sided (Section 3): hold the lower bracket $s_t \geq \frac{1}{2}\sigma_t$ everywhere, and keep the upward variation W_s a small multiple of V_σ^+ so that the $W_s^{1-1/p}$ factor of Theorem 3.3 is a constant, not a rate. A single-timescale trailing median cannot do both: a fast window tracks the scale but *churns* on heavy-tailed spikes ($W_s = 62 V_\sigma^+$), while a slow window is smooth but *lags* regime onsets and dips below the lower bracket. A two-timescale “peak-hold” envelope—react up on a short robust median, release down only geometrically—resolves the tension: it holds the stability-critical lower bracket at 100% of rounds while attaining the *smallest* upward variation, $W_s \approx 6.7 V_\sigma^+$ (Figure 6, Table 2). Its only cost is a benign 24% over-estimation of the scale (mild underfit, never divergence). Thus Assumption 3.2 is dischargeable in the direction that matters, and the scale drift enters the final bound only through the factor $\sqrt{W_s}$ —a modest constant here, since $W_s \approx 6.7 V_\sigma^+$ is a small multiple of the true drift. Section A.7 lifts this from an empirical observation to a proposition (Proposition A.5): under a piecewise-slowly-varying drift model the lower bracket holds off an $O(NW)$ -round set, reducing Assumption 3.2 to robust concentration on a near-stationary window.

Table 2. Scale trackers on the real gradient-scale process. The two-timescale envelope holds the lower bracket at 100% (stability) with the smallest upward variation W_s (regret constant). V_σ^+ is the true upward scale drift.

Tracker	Lower bracket	W_s	W_s/V_σ^+
Trailing median, fast ($W=64$)	0.93	38.6k	$62\times$
Trailing median, slow ($W=500$)	0.99	5.0k	$8.1\times$
Two-timescale envelope (ours)	1.00	4.2k	$6.7\times$

5. Related Work

Robust estimation and clipping. That the empirical mean fails under heavy tails, and robust estimators (Catoni, median-of-means, trimmed mean) achieve sub-Gaussian deviations under a finite variance, is classical (Catoni, 2012; Devroye et al., 2016; Lugosi & Mendelson, 2019); Bubeck et al. (2013) import this into sequential decision-making. In optimization, gradient clipping yields high-probability convergence under heavy-tailed noise (Zhang et al., 2020; Gorbunov et al., 2020; Nguyen et al., 2023). Our normalization is a streaming, scale-decoupled variant whose value is specifically stability under a *drifting* scale.

Heavy-tailed online learning. Zhang et al. (2026) show plain OGD is optimal *in expectation* on a bounded domain, and Zhang & Cutkosky (2022) give parameter-free *high-probability* static regret $\tilde{O}(\|u\|T^{1/p})$ over unbounded domains, whose central obstacle—exponentially large iterates—our cap sidesteps. Both are static; the dynamic, nonstationary regime we target sits outside them, and our lower bound shows why it is genuinely harder.




Figure 3. **Main result** (continuous multi-regime stream, *causal* standardization). (a) Weighted R^2 vs. learning rate: SN-OMD climbs monotonically to ≈ 0.16 and never diverges, while OGD and both clippers cliff above $\text{lr} = 10^{-2}$ (off-scale markers). (b) Peak rolling loss: SN-OMD stays bounded (≈ 5) everywhere; the others explode.




Figure 4. **The cap has an interior optimum** (synthetic, $p = 1.5$). Steady-state excess loss vs. M , each cap at its own tuned step: the finite- M optimum beats both the normalized-GD end ($M \rightarrow 0$) and the uncapped scale-adaptive-OGD end ($M \rightarrow \infty$).

Dynamic, scale-free, and adaptive OL. Dynamic regret against a moving comparator goes back to Zinkevich (2003); scale-free methods adapt to unknown gradient scale (Orabona & Pál, 2018); strongly-adaptive and parameter-free wrappers adapt to unknown interval/regime structure (Daniely et al., 2015; Cutkosky, 2020). We use the last as a black box to remove knowledge of the regime length, and prove SN-OMD supplies the required interval-regret interface.

Adaptive optimizers. Coordinate-wise adaptive methods (AdaGrad, RMSProp, Adam) also rescale the gradient by a running second-moment estimate and superficially resemble our normalization, so a natural question is whether they already provide robustness to a drifting scale. The difference is decisive under a heavy tail: their preconditioner divides




Figure 5. **Known-tail synthetic.** Cumulative regret vs. T at $p \in \{1.3, 1.5, 2.0\}$. OGD's exponent climbs past 1 (divergent) as the tail heavies; SN-OMD stays sublinear and below the $T^{1/p}$ reference (dashed).

by $\sqrt{\text{EMA}(g_t^2)}$, which *itself* spikes with a heavy-tailed gradient (the same round enters numerator and denominator), and there is no scale-free cap, so a single extreme gradient still yields an unbounded step. Our unconditional stability (Theorem 3.1) comes precisely from capping the *normalized* gradient at a constant M ; adaptive rescaling alone does not prevent the divergence documented in Section 4. An empirical head-to-head against the Adam family is a natural addition.

6. Conclusion

On real streams the difficulty is dominated by a strongly *non-stationary* gradient scale—which alone makes the pooled gradient look heavy-tailed, though the scale-normalized gradient is much lighter—and a budget heuristic suggests that this drift, against a moving comparator, is what pushes *dynamic* regret well above $\sqrt{T}$. The right response is to decouple the update from the tail scale: SN-OMD normalizes




Figure 6. Discharging Assumption 3.2. A two-timescale “peak-hold” scale tracker on the real gradient-scale process (a) holds the stability-critical lower bracket $s_t \geq \frac{1}{2}\sigma_t$ at every round, while (b) attaining the smallest upward variation W_s/V_σ^+ . A fast trailing median tracks the scale but churns (62×); a slow one is smooth but lags regime onsets and violates the lower bracket.

the gradient by a tracked scale and caps it at a scale-free constant, so its iterates provably stay bounded (for any learning rate) and it attains the per-regime worst-tail rate $T^{1/p}$, with the scale drift entering only as a $\sqrt{W_s}$ constant. Empirically it stays bounded across a $1000\times$ learning-rate range where OGD and scale-dependent clipping diverge.

Limitations. SN-OMD is not fully tuning-free: its best learning rate differs across settings (~ 1 on the continuous stream vs. ~ 0.5 on restarts), and its cap M , decay, and winsorization are untuned defaults—though its stable range is $\sim 1000\times$ wider than clipping’s and it degrades gracefully. The proven $W_s = O(V_\sigma^+)$ guarantee is for the two-timescale envelope tracker, whereas the empirically strongest default is a winsorized EMA; closing that gap with a single tracker that is both provably bounded-variation and empirically best is open, as is making the envelope’s release rate ρ adaptive to an unknown regime length. Consequently the dynamic-regret bound is fully proven for the envelope variant under Assumptions 2.1 and 3.2, while the deployed EMA variant’s high-probability dynamic regret under drifting (σ_t, p_t) is supported empirically. On the empirical side, three gaps remain. Our measurements come from one (large) market dataset; confirming the heavy-tail $\times$ nonstationarity phenomenology on a second domain (e.g. a synthetic drifting-tail stream, or RL/language-model gradients) would substantiate the generality claim. Our baselines are the theoretically matched methods—OGD, clipping, robust-mean estimators—and do not yet include the Adam family of coordinate-wise adaptive optimizers, whose empirical comparison we argue for but leave to future work (Section 5). And the evaluation centers on final weighted R^2 and iterate stability rather than mechanism-level diagnostics (regret trajectories, effective step-size dynamics, tracker error, seed

variance), which would more directly exhibit the stabilization mechanism. Open directions include per-coordinate normalization and tighter dynamic constants.

Software and Data

An implementation (the `ScaleNormalizedOGD` learner and scale trackers), configs, and scripts reproducing every figure and number will be released; the data is the public Jane Street Kaggle competition.

Impact Statement

This paper presents work whose goal is to advance the field of machine learning. There are many potential societal consequences of our work, none of which we feel must be specifically highlighted here.

References

- Bubeck, S., Cesa-Bianchi, N., and Lugosi, G. Bandits with heavy tail. *IEEE Transactions on Information Theory*, 59(11):7711–7717, 2013.
- Catoni, O. Challenging the empirical mean and empirical variance: A deviation study. *Annales de l’Institut Henri Poincaré, Probabilités et Statistiques*, 48(4):1148–1185, 2012.
- Cutkosky, A. Parameter-free, dynamic, and strongly-adaptive online learning. In *International Conference on Machine Learning (ICML)*, pp. 2250–2259, 2020.
- Daniely, A., Gonen, A., and Shalev-Shwartz, S. Strongly

adaptive online learning. In *International Conference on Machine Learning (ICML)*, pp. 1405–1411, 2015.

Devroye, L., Lerasle, M., Lugosi, G., and Oliveira, R. I. Sub-Gaussian mean estimators. *The Annals of Statistics*, 44(6):2695–2725, 2016.

Gorbunov, E., Danilova, M., and Gasnikov, A. Stochastic optimization with heavy-tailed noise via accelerated gradient clipping. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

Jane Street. Jane Street real-time market data forecasting. Kaggle competition, 2024. <https://www.kaggle.com/competitions/jane-street-real-time-market-data-forecasting>.

Lugosi, G. and Mendelson, S. Mean estimation and regression under heavy-tailed distributions: A survey. *Foundations of Computational Mathematics*, 19(5):1145–1190, 2019.

Nguyen, T. D., Ene, A., and Nguyen, H. L. High probability convergence of clipped-SGD under heavy-tailed noise. *arXiv preprint arXiv:2302.05437*, 2023.

Orabona, F. and Pál, D. Scale-free online learning. *Theoretical Computer Science*, 716:50–69, 2018.

Zhang, J. and Cutkosky, A. Parameter-free regret in high probability with heavy tails. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022. arXiv:2210.14355.

Zhang, J., He, T., Sra, S., and Jadbabaie, A. Why gradient clipping accelerates training: A theoretical justification for adaptivity. In *International Conference on Learning Representations (ICLR)*, 2020.

Zhang, Z., Zhang, M., and Zhou, Z. Online convex optimization with heavy tails: Old algorithms, new regrets, and applications. In *Algorithmic Learning Theory (ALT)*, 2026. arXiv:2508.07473.

Zinkevich, M. Online convex programming and generalized infinitesimal gradient ascent. In *International Conference on Machine Learning (ICML)*, pp. 928–936, 2003.

A. Proofs

Throughout, $v_t = g_t/s_t$, $\hat{g}_t = \text{clip}(v_t, M)$ with $M \geq 1$, and $p = \inf_t p_t$. We use the filtration $(\mathcal{F}_t)$ with $\mathbb{E}[g_t | \mathcal{F}_{t-1}] = \bar{g}_t$ and s_t predictable ($\mathcal{F}_{t-1}$ -measurable).

A.1. Stability (Theorem 3.1)

Since $\hat{g}_t = v_t \min\{1, M/\|v_t\|\}$, $\|\hat{g}_t\| \leq M$ for every realization. On a bounded domain $\Pi_{\mathcal{W}}$ keeps $\|w_t\| \leq D$. On $\mathbb{R}^d$, $\|w_{t+1} - w_t\| \leq (\eta/\sqrt{t})M$, so $\|w_t\| \leq \|w_1\| + M\eta \sum_{k < t} k^{-1/2} \leq \|w_1\| + 2M\eta\sqrt{t}$. No moment assumption is used. $\square$

A.2. Clipped moment and bias

Lemma A.1. For $v \in \mathbb{R}^d$, $M > 0$, $p \in (1, 2]$: $\|\text{clip}(v, M)\|^2 \leq M^{2-p}\|v\|^p$ and $\|v - \text{clip}(v, M)\| \leq \|v\|^p/M^{p-1}$.

Proof. If $\|v\| \leq M$: $\text{clip}(v, M) = v$ and $\|v\|^2 = \|v\|^p\|v\|^{2-p} \leq \|v\|^p M^{2-p}$, while $v - \text{clip}(v, M) = 0$. If $\|v\| > M$: $\|\text{clip}(v, M)\|^2 = M^2 = M^p M^{2-p} \leq \|v\|^p M^{2-p}$, and $\|v - \text{clip}(v, M)\| = \|v\| - M \leq \|v\| = \|v\|^p\|v\|^{1-p} < \|v\|^p M^{1-p}$. $\square$

Lemma A.2 (Moment and bias under Assumption 3.2). There is a constant κ with $\mathbb{E}[\|\hat{g}_t\|^2 | \mathcal{F}_{t-1}] \leq M^{2-p}\kappa$, and the bias $b_t = \mathbb{E}[v_t - \hat{g}_t | \mathcal{F}_{t-1}]$ obeys $\|b_t\| \leq \kappa/M^{p-1}$.

Proof. Take conditional expectations in Lemma A.1 with $v = v_t$ and use $\mathbb{E}[\|g_t\|^{p_t} | \mathcal{F}_{t-1}] \leq \kappa_0 \sigma_t^{p_t}$ and $s_t \geq c_1 \sigma_t$: $\mathbb{E}\|\hat{g}_t\|^2 \leq M^{2-p}\mathbb{E}\|v_t\|^p \leq M^{2-p}\kappa_0 c_1^{-p}$, and $\|b_t\| \leq \mathbb{E}\|v_t\|^{p_t}/M^{p_t-1} \leq \kappa_0 c_1^{-p_t}/M^{p-1}$ (using $M \geq 1$, $p_t \geq p$). Set $\kappa = \kappa_0 c_1^{-2}$. $\square$

Because $M \geq 1$ and $p_t \geq p$, all uses of M^{2-p_t} , $M^{-(p_t-1)}$ above are bounded by the worst-case M^{2-p} , $M^{-(p-1)}$; the single- p bounds are thus valid for drifting p_t .

A.3. High probability

Lemma A.3 (Freedman step). Let (X_t) be a martingale difference sequence with $|X_t| \leq b$ and $\sum_{t \leq T} \mathbb{E}[X_t^2 | \mathcal{F}_{t-1}] \leq V$ almost surely. Then w.p. $\geq 1 - \delta$, $\sum_{t \leq T} X_t \leq \sqrt{2V \log(1/\delta)} + \frac{2b}{3} \log(1/\delta)$. In particular, for feedback Z_t with $\|Z_t\| \leq B$ and $X_t = \langle Z_t - \mathbb{E}[Z_t | \mathcal{F}_{t-1}], w_t - u_t \rangle$ (u_t deterministic, $\|w_t - u_t\| \leq D$), one may take $b = 2BD$ and $V = D^2 \sum_t \mathbb{E}[\|Z_t\|^2 | \mathcal{F}_{t-1}]$.

Proof. (X_t) is a martingale difference sequence ($\mathbb{E}[X_t | \mathcal{F}_{t-1}] = 0$). Freedman's inequality gives $\Pr(\sum_t X_t \geq \lambda) \leq \exp(-\lambda^2/(2V + 2b\lambda/3))$; inverting the right-hand side to δ and using $\sqrt{x+y} \leq \sqrt{x} + \sqrt{y}$ yields the bound. For the instantiation, $|X_t| \leq \|Z_t - \mathbb{E}[Z_t | \mathcal{F}_{t-1}]\| \|w_t - u_t\| \leq 2BD$ and $\mathbb{E}[X_t^2 | \mathcal{F}_{t-1}] \leq D^2 \mathbb{E}[\|Z_t\|^2 | \mathcal{F}_{t-1}]$ since $\mathbb{E}\|Z - \mathbb{E}Z\|^2 \leq \mathbb{E}\|Z\|^2$. $\square$

A.4. Varying-step telescope

Lemma A.4. SN-OMD is OGD with step $\eta_t = \eta/s_t$ on the clipped raw gradient. For any comparator sequence $u_{1:T}$, writing $W_s = s_1 + \sum_t (s_t - s_{t-1})_+$ and $P_T^s = \sum_t s_t \|u_t - u_{t-1}\|$,

$$\sum_{t=1}^T s_t (\|w_t - u_t\|^2 - \|w_{t+1} - u_t\|^2) \leq D^2 W_s + 2D P_T^s.$$

Proof. Reindexing the sum and splitting each bracket into a comparator-move and a scale-change part,

$$s_t \|w_t - u_t\|^2 - s_{t-1} \|w_t - u_{t-1}\|^2 = s_t (\|w_t - u_t\|^2 - \|w_t - u_{t-1}\|^2) + (s_t - s_{t-1}) \|w_t - u_{t-1}\|^2,$$

which is at most $2Ds_t \|u_t - u_{t-1}\| + (s_t - s_{t-1})_+ D^2$ by the reverse triangle inequality and $\|\cdot\| \leq D$. Summing and dropping the nonpositive boundary term gives the bound. $\square$

A.5. Dynamic regret (Theorem 3.3)

Roadmap and where each assumption enters. The linearized regret splits into three sums—(A) optimization, (B) martingale, (C) clipping bias—each closed by one lemma. (A) uses the varying-step telescope (Lemma A.4) for the comparator/ scale-drift terms and the clipped second moment (Lemmas A.1 and A.2) for the curvature; (B) uses the Freedman step (Lemma A.3); (C) uses the clipping-bias bound (Lemma A.2). The two assumptions enter in exactly two isolated places: Assumption 2.1 (finite p -th moment) only through Lemma A.2, which converts $\mathbb{E}\|g_t\|^{p_t} \leq \sigma_t^{p_t}$ into the conditional second-moment and bias bounds of (A)–(C); and Assumption 3.2 (the lower bracket $s_t \geq c_1 \sigma_t$) only in the curvature of (A)/(B) and the bias of (C)—*nowhere else*, which is precisely why over-estimating the scale is harmless. The dependency is $\text{Lemma A.1} \rightarrow \text{Lemma A.2} \rightarrow \{(A), (B), (C)\}$, with Lemma A.4 feeding (A) and Lemma A.3 feeding (B); the final display balances (A)’s curvature against (C)’s bias through the choice of η and M .

Write $d_t := \text{clip}(g_t, Ms_t) = s_t \hat{g}_t$ for the clipped *raw* gradient, so that SN-OMD is exactly OGD with adaptive step $\eta_t = \eta/s_t$ on d_t (Lemma A.4). By convexity $R_T(u_{1:T}) \leq \sum_t \langle \bar{g}_t, w_t - u_t \rangle$, and we split the true subgradient as $\bar{g}_t = m'_t + \beta_t$, where $m'_t = \mathbb{E}[d_t \mid \mathcal{F}_{t-1}]$ and $\beta_t = \bar{g}_t - m'_t = \mathbb{E}[(g_t - d_t) \mid \mathcal{F}_{t-1}]$ is the (raw) clipping bias. Then

$$\sum_t \langle \bar{g}_t, w_t - u_t \rangle = \underbrace{\sum_t \langle d_t, w_t - u_t \rangle}_{\text{(A) optimization}} + \underbrace{\sum_t \langle m'_t - d_t, w_t - u_t \rangle}_{\text{(B) martingale}} + \underbrace{\sum_t \langle \beta_t, w_t - u_t \rangle}_{\text{(C) bias}}.$$

(A) The OGD one-step inequality with step η_t , summed and combined with Lemma A.4, gives $\sum_t \langle d_t, w_t - u_t \rangle \leq \frac{1}{2\eta} (D^2 W_s + 2DP_T^s) + \frac{\eta}{2} \sum_t \|d_t\|^2 / s_t$. Since $\|d_t\|^2 = \|\text{clip}(g_t, Ms_t)\|^2 \leq (Ms_t)^{2-p} \|g_t\|^p$ (Lemma A.1) and $s_t \geq c_1 \sigma_t$ (Assumption 3.2), the curvature sum has conditional mean $\leq M^{2-p} \kappa' \sum_t \sigma_t = M^{2-p} \kappa' S_1$, with $S_1 = \sum_t \sigma_t$. (B) The increments $\langle m'_t - d_t, w_t - u_t \rangle$ form a martingale difference with $\|d_t\| \leq Ms_t$ and $\mathbb{E}[\|d_t\|^2 \mid \mathcal{F}_{t-1}] \leq M^{2-p} \kappa' \sigma_t^2$; Lemma A.3 bounds this sum by $\tilde{O}(DM^{1-p/2} \sqrt{\sup_t \sigma_t \cdot S_1})$, the same order as (A). (C) $\|\beta_t\| \leq \mathbb{E}[\|g_t\|^p \mid \mathcal{F}_{t-1}] / (Ms_t)^{p-1} \leq \kappa' \sigma_t / M^{p-1}$ (Lemma A.2), so (C) $\leq D\kappa' S_1 / M^{p-1}$. Choosing $\eta = \Theta(\sqrt{(D^2 W_s + DP_T^s) / (M^{2-p} S_1)})$ and $M = \Theta(T^{1/p})$ balances the curvature term of (A) against the bias (C) and yields

$$R_T(u_{1:T}) = O(\sqrt{(D^2 W_s + DP_T^s) M^{2-p} S_1} + DS_1 / M^{p-1}) = \tilde{O}((\sup_t \sigma_t)(D + \sqrt{DP_T}) T^{1/p} \log(1/\delta)),$$

where the last step substitutes $M = \Theta(T^{1/p})$, $S_1 \leq T \sup_t \sigma_t$, and $W_s, P_T^s \lesssim \sup_t \sigma_t \cdot (1, P_T)$; equivalently the first term is $(D\sqrt{W_s} + \sqrt{DP_T^s}) \sqrt{S_1} T^{(2-p)/(2p)}$. The scale drift enters as the factor $\sqrt{W_s}$, not as a rate. The stationary case $\sigma_t \equiv \sigma$, $p = 2$ recovers $\sigma(D + \sqrt{DP_T}) \sqrt{T}$, the classical dynamic-OGD rate. $\square$

A.6. Reductions

Unknown p : $M = \Theta(T^{1/p})$ needs p ; a doubling schedule over M , using that the bound is unimodal in $\log M$ with a flat optimum, removes this at a polylog(T) cost (Zhang et al., 2026). *Unknown $\|u\|$ / unbounded domain:* Theorem 3.1 gives bounded feedback $\|\hat{g}_t\| \leq M$, exactly the input required by the parameter-free potential of Zhang & Cutkosky (2022); substituting our $\hat{g}_t$ yields the unbounded-domain analogue with D replaced by $\|u\|$. *Unknown regime length:* Lemma A.4 plus the interval version of Lemma A.3 give SN-OMD an interval-regret guarantee on every $[a, b]$; a strongly-adaptive wrapper (Daniely et al., 2015; Cutkosky, 2020) then recovers the dynamic bound without knowing the regime structure, at a polylog overhead. A per-regime transition cost $O(DN \text{polylog } T^{1/p})$ is incurred while the scale tracker re-concentrates after each of the N regime changes; it is dominated when regimes are macroscopic.

A.7. Discharging Assumption 3.2: a bracket guarantee

Assumption 3.2 posits a predictable scale that brackets σ_t . We show it is not an oracle: under a mild model of the drift, the two-timescale envelope of Section 4.3 *provably* meets the essential lower bracket off a small exceptional set, so the burden moves from “an accurate scale estimate exists” to “the drift is piecewise slowly varying,” and the estimate is then *derived* from standard robust concentration.

Drift model. Partition $[T]$ into N regimes by change points $t_0 < \dots < t_N$. *Slow within-regime drift:* inside a regime the scale varies slowly, $\sigma_t / \sigma_{t-1} \in [1 - \gamma, 1 + \gamma]$ with $\gamma W \leq c_0$ for the window W . *Arbitrary jumps:* at the N change points σ_t may jump by any factor. This is the piecewise-slowly-varying picture the measured gradient scale fits ($\sim 30\%$ within-day drift, occasional $3\times$ overnight jumps; Section 4).

Tracker. The envelope $s_t = \max\{c\hat{m}_t, (1 - \rho)s_{t-1}\}$, where $\hat{m}_t = \text{median}(\|g_{t-W}\|, \dots, \|g_{t-1}\|)$ is the predictable short-window median and $\rho \asymp 1/L$ with $L = \min_r L_r$.

Proposition A.5 (Bracket guarantee). *Under Assumption 2.1 and the drift model above, for a large enough constant c , with probability $\geq 1 - \delta$: (i, lower bracket) $s_t \geq c_1\sigma_t$ for every t outside at most $O(W)$ rounds after each jump—an exceptional set of size $O(NW)$; and (ii, bounded variation) $W_s = O(\sup_t \sigma_t + \rho \sum_t \sigma_t) = O(V_\sigma^+)$. Consequently Theorem 3.3 holds with one added term $O(DNW \sup_t \sigma_t)$, dominated whenever $NW \ll T^{1/p}$ (macroscopic regimes).*

Proof sketch. **Window concentration.** On a window lying inside one regime, the population median of $\|g_t\|$ is a constant multiple of σ_t (regular variation), and the empirical median concentrates around it by order-statistic bounds (needing essentially no moment); the slow drift shifts the constants by $O(\gamma W) = O(1)$. Hence $\hat{m}_t \in [a\sigma_t, b\sigma_t]$ w.h.p., and $c \geq c_1/a$ gives $c\hat{m}_t \geq c_1\sigma_t$, so the max meets the lower bracket whenever the window lies in one regime. **Jumps.** After an upward jump the window still holds pre-jump norms for $\leq W$ rounds; during that lapse the median term may under-estimate, giving the $O(NW)$ exceptional set. A downward jump only makes s_t conservatively high (the release lets it fall at rate ρ), never violating the lower bracket. **Variation.** Upward moves of s_t come only from $c\hat{m}_t$ rising; between jumps these telescope to $\sup_t \sigma_t$, and the geometric release contributes the churn $\rho \sum_t s_t$, so $W_s \leq \sup_t \sigma_t + \rho \sum_t s_t$, which is $O(V_\sigma^+)$ for $\rho \asymp 1/L$ (Remark 3.5). **Exceptional rounds.** On the $O(NW)$ lapse rounds the update is still bounded ($\|\hat{g}_t\| \leq M$, Theorem 3.1) but the clipping bias (C) is uncontrolled; charging it trivially at $O(D \sup_t \sigma_t)$ per round gives the added $O(DNW \sup_t \sigma_t)$. $\square$

The residual—the post-jump lapse—is an additive term already of the form isolated in the Reductions above, and making ρ adaptive to an unknown L is exactly the strongly-adaptive wrapper of that subsection.

A.8. The deployed EMA vs. the analyzed envelope tracker

The regret analysis discharges Assumption 3.2 with the two-timescale “peak-hold” envelope of Section 4.3, whereas the empirically strongest configuration uses a winsorized exponential moving average (EMA). Because a reviewer may reasonably ask whether the analyzed object and the deployed object are the same, we state precisely what does and does not depend on the choice of tracker.

Stability is tracker-independent. Theorem 3.1 uses only $\|\hat{g}_t\| \leq M$, which the cap enforces for any predictable $s_t > 0$. Non-divergence therefore holds verbatim for the EMA, for the envelope, and for any other tracker; the gap discussed here is never about stability, only about the leading regret constant.

Only the regret constant depends on the tracker. For any predictable s_t meeting the essential lower bracket $s_t \geq c_1\sigma_t$, the rate is $T^{1/p}$ (Theorem 3.3); trackers differ only through the leading factor $W_s^{1-1/p}$. The envelope is the tracker for which we can prove both the lower bracket and $W_s = O(V_\sigma^+)$ (Proposition A.5)—reacting up on a short robust median and releasing down only geometrically keeps every upward move a genuine rise (Table 2: 100% lower bracket, $W_s = 6.7 V_\sigma^+$). A fast, *unwinsorized* EMA can in principle churn on heavy-tailed $\|g_t\|$ spikes (Table 2: $W_s = 62 V_\sigma^+$); winsorizing the EMA at $k \cdot$ median caps each upward step and is what keeps the deployed default’s W_s finite in practice, though we do not give it a closed-form bound.

Why deploy the EMA nonetheless. It is simpler, standard, and empirically more accurate—weighted $R^2 \approx 0.24$ versus ≈ 0.19 for the envelope on the continuous stream—because the envelope’s benign over-estimation of the scale shrinks its steps η/s_t into a mild underfit, which is precisely the price of its provable W_s . The two trackers are thus endpoints of a single accuracy/guarantee trade-off, and *both* are bounded; deployment can pick either without risking divergence.

Closing the gap. Either run the envelope in deployment, accepting the small accuracy cost, or prove a W_s bound for the winsorized EMA under a heavy-tailed scale process: the winsorization already supplies the per-step cap the envelope gets from its median, so what remains is to control the geometric-decay churn term $\rho \sum_t s_t$, exactly as in Remark 3.5, with ρ the EMA rate. The one residual modeling assumption— ρ matched to an unknown regime length—is removed by the strongly-adaptive wrapper of the preceding subsection.